

# ClaimSight

## Refunds with receipts.

An after-sales teammate that watches the customer's evidence video, applies your refund policy, executes the refund itself, and catches returns fraud by matching footage across accounts.

CLAIM C-1042 · ORDER A1042

# \$24.00 refunded

Auto-approved · 18.3 s

Evidence chip on rim @ 0:03

Policy P1 P2 P4 ✓

Transaction txn\_8f3a...

EdgeOne Makers Memories.ai AgentX WorkBuddy

### FRAUD TWIN DETECTED

Same footage as C-1042 from a different account (similarity 0.91). Claim frozen, manager pinged in Slack.

# Every refund still needs a human to squint at a photo.

**\$101B**

## lost to return fraud

US retailers, 2023. \$13.70 of every \$100 returned is fraud or abuse. (NRF and Appriss Retail, 2023 Consumer Returns in the Retail Industry)

**0**

## people watching the evidence

Customers send photos and videos. Agents skim them, guess, and either overpay or escalate everything to a queue.

**3**

## claims, one cracked mug

The same damaged item gets filmed for three claims from three "customers". No refund queue can see across accounts.

# An after-sales teammate that finishes the job.

1

## Sees the evidence

Memories.ai indexes the customer's 10-second clip: "white ceramic mug, chip on the rim at 0:03". Captions, frames and a summary become the claim's evidence record.

2

## Applies your policy

Orders and policy clauses P1–P6 live as structured JSON in Makers storage. No vector database. Every decision cites the clauses it used.

3

## Executes the transaction

Refund or replacement is written to the ledger by a cloud function that re-checks the limits itself. Humans are pulled in only above \$75 or on a fraud signal. The receipt is a Blender-rendered 3D twin of the damage.



Damage Twin receipt: rendered in Blender where Memories.ai saw the damage

# One loop, from evidence to ledger entry.

## DATA

### Video + order

A 10-second clip.  
Orders and policy as  
JSON in Makers  
storage.

## INTELLIGENCE

### See the damage

Memories.ai captions  
the clip and image-  
searches every prior  
claim for a twin.

## EXPERIENCE

### Decision Card

Streaming chat shows  
the reasoning, cited  
clauses, evidence line  
and txn id.

## EXECUTION

### Policy, twice

The agent applies it;  
the refund service  
enforces it again.  
Over limit? A human  
in Slack.

## TRANSACTION

### Refund issued

Ledger entry, decision  
record, AgentX trace,  
dashboard row.

Most teams stop at a chat answer. ClaimSight reaches the last layer: a transaction with a receipt.

# A chipped mug. A refund. A fraud catch.

## 0:00 Claim

"My mug arrived chipped, order A1042." Evidence clip attached. The trace lights up: order, policy, evidence, fraud, execute, record.

## 0:30 Decision

Damage seen at 0:02. Window checked (P1), damage visible (P2), replacement-first rule applied (P4). \$24 is under the \$75 limit: refund executed, transaction id on the card, and a 3D Damage Twin renders into the card 25 s later.

## 1:00 Twin

Same footage from a different account. Image search finds it at 0.93 similarity. Claim frozen, escalated to Slack, denied by the manager in the Refund Desk.

The screenshot displays the ClaimSight interface, which is a dark-themed dashboard for managing claims. The top navigation bar includes 'Chat', 'Refund Desk', and 'Lab' tabs, along with a 'DEMO DATA' button and language settings (ZH, EN). The main content area is divided into several sections:

- CUSTOMER:** Alice Moreno (e\_alice)
- EVIDENCE:** White ceramic mug, chip on the rim at 0:03
- POLICY CLAUSES:** P2, P4
- TRANSACTION:** ✓ txn\_mu8j2c7a
- FRAUD CHECK:** No matching footage in other claims

The central part of the interface features a 3D rendering of a white ceramic mug with a chip on the rim, labeled 'chipped mug, broken rim'. Below the rendering, a card displays 'CLAIMSIGHT - DAMAGE TWIN C-A1042-81D3 Alice Moreno - Order A1042' and 'Refund issued - \$24.00'.

On the right side, the **Decision trace** section shows the decision was made in 33 ms by the POLICY ENGINE (C-A1042-81D3). Below this, the **EVIDENCE** section shows a clip titled 'CLIP vid\_stub\_mug\_alice' with a thumbnail of the mug. The **PIPELINE** section lists the following steps:


- Order (lookup\_order) - <1 ms
- Policy (get\_policy) - <1 ms
- Evidence (inspect\_evidence) - <1 ms
- Fraud (fraud\_check) - <1 ms
- Execute (execute\_refund) - <1 ms
- Record (done) - <1 ms

The bottom of the interface shows a list of scenarios: A1042 Chipped mug, A1043 Same mug, other account, and A1050 Cracked headphones. Below this, there is a section for 'ORDER A1042' with details 'White ceramic mug - \$24.00 - delivered 3d ago' and buttons for 'Attach evidence' and 'Demo clip'. A text input field for 'Describe the problem with your order...' and a 'File claim' button are also present.

# Agent runtime, cloud functions and storage on EdgeOne Makers; co-hosts do the seeing, gating and reviewing.

## ClaimSight architecture

Evidence video → policy decision → executed transaction, with every step traced and gated




# The same mug, a different account.

Every evidence clip is indexed into the claims collection. Before paying, ClaimSight takes a frame from the new clip and runs an image search across every prior claim. The detector was crash-tested with a Blender-rendered synthetic evidence set.

# 0.93

similarity to the claim on order A1042, filed by another customer

- Policy P5 freezes the claim instead of paying it.
- The manager gets a Slack ping through the WorkBuddy Refund Desk skill.
- One click in the desk: approve or deny, with a note that lands in the audit record.



Synthetic Evidence Lab · detector: local-baseline

## 12

synthetic clips  
rendered in Blender

## 75%

twin recall at the  
shipped threshold

## 18%

false-positive rate

## 0.76

similarity threshold  
from data

# Shipped like software.

30

golden claims

clean, over-limit, out-of-window, replacement-first, fraud twin, non-returnable

3

online monitors

PII, secrets and prompt-injection scorers on every reply

7.5

gate floor out of 10

plus a no-regression check against the previous run; CI fails the build below it

1

trace per claim

every tool call, latency and decision in AgentX, linked to the eval run

The screenshot shows the 'AGENTX Observe - Live Traces' interface. On the left is a sidebar with navigation options: Overview, Observe (selected), Monitor, Review, Scorers, Topics, Evaluate, Insights, Playground, MANAGE (Datasets, Prompts, Tools & MCPs, Rules, Settings), AUTOMATIONS (CI Gates, Agent Skill), and ENTERPRISE (SSO, private deployment, high volume, Contact us). The main panel displays a table of live traces. The table has columns: TIME, TRACE, SOURCE, TOOL QUALITY, SCORE, IN TOKENS, OUT TOKENS, and LATENCY. It shows 10 rows of traces, all for 'agent claimsight.claim' at '6:20 AM', with a 'pass' status and a score of 4 or 5. The bottom of the table indicates 'Rows per page: 25 50 100' and '1-50 of 370'.

| TIME    | TRACE                  | SOURCE                 | TOOL QUALITY | SCORE   | IN TOKENS | OUT TOKENS | LATENCY |
|---------|------------------------|------------------------|--------------|---------|-----------|------------|---------|
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 4       | 70.0000 | -         | -          | 1ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 4       | 70.0000 | -         | -          | 2ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 4       | 70.0000 | -         | -          | 1ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 5       | 70.0000 | -         | -          | 4ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 5       | 70.0000 | -         | -          | 3ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 5       | 70.0000 | -         | -          | 2ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 5       | 70.0000 | -         | -          | 3ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 5       | 70.0000 | -         | -          | 3ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 5       | 70.0000 | -         | -          | 2ms     |
| 6:20 AM | agent claimsight.claim | agent claimsight.claim | pass 5       | 70.0000 | -         | -          | 3ms     |

AgentX Observe: one trace per ClaimSight claim, straight from the running backend.

**Limits live in the refund service, not the prompt:** the function re-checks amount, window and approval header, so the model cannot exceed policy even when asked to. The agent has no shell, no browser and no keys. If the gateway blinks, a deterministic policy engine runs the same steps and produces the same card.

# Every sponsor tool does real work.

## EdgeOne Makers

Agent runtime (session mode), custom MCP tools, Blob storage, AI Gateway models, streaming UI, tracing, one-command deploy.

## Memories.ai

Video Datalake: upload, index, summary and captions, frame extraction, image search across the claims collection for the fraud twin.

## AgentX

Self-hosted trace-eval: golden dataset, policy judge, deploy gate with no-regression check, online PII and secrets monitors.

## WorkBuddy

Refund Desk skill: pulls escalations, summarizes evidence and fraud matches, takes approve or deny, posts to Slack; daily .docx report.

## VeloDB

claims\_events table and dashboard SQL: auto-approval by hour, refunded dollars, fraud clusters, p95 latency.

## AWS

S3 evidence archive on every upload (private, versioned, encrypted) and a CloudFormation stack for the EC2 host that runs AgentX and the live-evidence relay.

# Data → Intelligence → Experience → Execution → Transaction. **One loop.**

## Why it wins

Returns fraud is a \$101B problem with a named buyer on every marketplace. The demo reaches a transaction, and the fraud catch is a moment judges remember.

## Next 30 days

One pilot with a marketplace returns desk. Measure auto-approval rate, fraud catches and minutes saved per claim. AWS credits fund video volume, S3 and the AgentX host.

## Try it

Live demo: [demo URL]  
Repository:  
[github.com/vnmoorthy/claimsight](https://github.com/vnmoorthy/claimsight)  
Deck and storyboard in docs/

# Thank you.