

ClaimSight

Refunds with receipts.

An after-sales teammate that watches the customer's evidence video, applies your refund policy, executes the refund itself, and catches returns fraud by matching footage across accounts.

CLAIM C-1042 · ORDER A1042

\$24.00 refunded

Auto-approved · 18.3 s

Evidence chip on rim @ 0:03

Policy P1 P2 P4 ✓

Transaction txn_8f3a...

EdgeOne Makers Memories.ai AgentX WorkBuddy

FRAUD TWIN DETECTED

Same footage as C-1042 from a different account (similarity 0.91). Claim frozen, manager pinged in Slack.

Every refund still needs a human to squint at a photo.

\$101B

lost to return fraud

US retailers, 2023. \$13.70 of every \$100 returned is fraud or abuse. (NRF and Appriss Retail, 2023 Consumer Returns in the Retail Industry)

0

people watching the evidence

Customers send photos and videos. Agents skim them, guess, and either overpay or escalate everything to a queue.

3

claims, one cracked mug

The same damaged item gets filmed for three claims from three "customers". No refund queue can see across accounts.

An after-sales teammate that finishes the job.

1

Sees the evidence

Memories.ai indexes the customer's 10-second clip: "white ceramic mug, chip on the rim at 0:03". Captions, frames and a summary become the claim's evidence record.

2

Applies your policy

Orders and policy clauses P1–P6 live as structured JSON in Makers storage. No vector database. Every decision cites the clauses it used.

3

Executes the transaction

Refund or replacement is written to the ledger by a cloud function that re-checks the limits itself. Humans are pulled in only above \$75 or on a fraud signal. The receipt is a Blender-rendered 3D twin of the damage.



Damage Twin receipt: rendered in Blender where the evidence showed the damage

One loop, from evidence to ledger entry.

DATA

Video + order

A 10-second clip.
Orders and policy as
JSON in Makers
storage.

INTELLIGENCE

See the damage

The clip is captioned
and its frames are
searched against
every prior claim for a
twin.

EXPERIENCE

Decision Card

Streaming chat shows
the reasoning, cited
clauses, evidence line
and txn id.

EXECUTION

Policy, twice

The agent applies it;
the refund service
enforces it again.
Over limit? A human
in Slack.

TRANSACTION

Refund issued

Ledger entry, decision
record, AgentX trace,
dashboard row.

Most teams stop at a chat answer. ClaimSight reaches the last layer: a transaction with a receipt.

A chipped mug. A refund. A fraud catch.

0:00 Claim

"My mug arrived chipped, order A1042." Evidence clip attached. The trace lights up: order, policy, evidence, fraud, execute, record.

0:30 Decision

Damage seen at 0:02. Window checked (P1), damage visible (P2), replacement-first rule applied (P4). \$24 is under the \$75 limit: refund executed, transaction id on the card, and a 3D Damage Twin renders into the card 25 s later.

1:00 Twin

Same footage from a different account: matched at 0.93 similarity across prior claims. Claim frozen, escalated to Slack, denied by the manager in the Refund Desk.

The screenshot displays the ClaimSight interface, which is a dark-themed dashboard for managing claims. The top navigation bar includes tabs for Chat, Refund Desk, and Lab, along with a DEMO DATA button and language settings (ZH, EN). The main content area is divided into several sections:

- CUSTOMER:** Alice Moreno (e_alice)
- EVIDENCE:** White ceramic mug, chip on the rim at 0:03
- POLICY CLAUSES:** P2, P4
- TRANSACTION:** txn_mu8j2c7a
- FRAUD CHECK:** No matching footage in other claims

The central part of the interface features a 3D rendering of a white ceramic mug with a chip on the rim, labeled "chipped mug, broken rim". Below the rendering, a card displays "White ceramic mug, chip on the rim at 0:03" and "Refund issued - \$24.00".

On the right side, the **Decision trace** section shows a decision made in 33 ms by the POLICY ENGINE (C-A1042-8103). Below this, the **EVIDENCE** section displays a clip titled "CLIP vid_stub_mug_alice" with a thumbnail showing a white ceramic mug. The **PIPELINE** section lists the following steps:


- Order (lookup_order) - <1 ms
- Policy (get_policy) - <1 ms
- Evidence (inspect_evidence) - <1 ms
- Fraud (fraud_check) - <1 ms
- Execute (execute_refund) - <1 ms
- Record (done) - <1 ms

At the bottom, the **SCENARIOS** section lists A1042 (Chipped mug), A1043 (Same mug, other account), and A1050 (Cracked headphones). The **ORDER** A1042 section shows "White ceramic mug - \$24.00 - delivered 3d ago" and buttons for "Attach evidence" and "Demo clip". A text input field for "Describe the problem with your order..." and a "File claim" button are also present.

Agent runtime, cloud functions and storage on EdgeOne Makers; co-hosts do the seeing, gating and reviewing.

ClaimSight architecture

Evidence video → policy decision → executed transaction, with every step traced and gated




The same mug, a different account.

Every evidence clip is indexed. Before paying, ClaimSight takes a frame from the new clip and searches every prior claim for the same footage. The detector was crash-tested with a Blender-rendered synthetic evidence set.

0.93

similarity to the claim on order A1042, filed by another customer

- Policy P5 freezes the claim instead of paying it.
- The manager gets a Slack ping through the WorkBuddy Refund Desk skill.
- One click in the desk: approve or deny, with a note that lands in the audit record.



Synthetic Evidence Lab · detector: local-baseline

12

synthetic clips
rendered in Blender

75%

twin recall at the
shipped threshold

18%

false-positive rate

0.76

similarity threshold
from data

Shipped like software.

30

golden claims

clean, over-limit, out-of-window, replacement-first, fraud twin, non-returnable

3

online monitors

PII, secrets and prompt-injection scorers on every reply

7.5

gate floor out of 10

plus a no-regression check against the previous run; CI fails the build below it

1

trace per claim

every tool call, latency and decision in AgentX, linked to the eval run

The screenshot shows the 'AGENTX Observe - Live Traces' interface. It features a sidebar with navigation options like Overview, Monitor, Review, Scorer, Topics, Evaluate, Insights, and Playground. The main area displays a table of live traces. The table has columns for TIME, TRACE, SOURCE, TOOL QUALITY, SCORE, IN TOKENS, OUT TOKENS, and LATENCY. The traces are for 'agent claimsight.claim' and show a 'pass' status with a score of 4 or 5. The latency is consistently low, ranging from 1ms to 4ms. The interface also includes a search bar and a 'Framework' dropdown.

TIME	TRACE	SOURCE	TOOL QUALITY	SCORE	IN TOKENS	OUT TOKENS	LATENCY
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 4	70.0000	-	-	1ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 4	70.0000	-	-	2ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 4	70.0000	-	-	1ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	4ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	2ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	2ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms

AgentX Observe: one trace per ClaimSight claim, straight from the running backend.

Limits live in the refund service, not the prompt: the function re-checks amount, window and approval header, so the model cannot exceed policy even when asked to. The agent has no shell, no browser and no keys. If the gateway blinks, a deterministic policy engine runs the same steps and produces the same card.

Every sponsor tool does real work.

EdgeOne Makers

Agent runtime (session mode), custom MCP tools, Blob storage, AI Gateway models, streaming UI, tracing, one-command deploy.

Memories.ai

Video Datalake client: upload, index, captions, frames, image search across prior claims. Built and contract-tested; demo day ran on canned evidence.

AgentX

Self-hosted trace-eval: golden dataset, policy judge, deploy gate with no-regression check, online PII and secrets monitors.

WorkBuddy

Refund Desk skill: pulls escalations, summarizes evidence and fraud matches, takes approve or deny, posts to Slack; daily .docx report.

VeloDB

claims_events table and dashboard SQL: auto-approval by hour, refunded dollars, fraud clusters, p95 latency.

AWS

S3 evidence archive on every upload (private, versioned, encrypted) and a CloudFormation stack for the EC2 host that runs AgentX and the live-evidence relay.

Data → Intelligence → Experience → Execution → Transaction. **One loop.**

Why it wins

Returns fraud is a \$101B problem with a named buyer on every marketplace. The demo reaches a transaction, and the fraud catch is a moment judges remember.

Next 30 days

One pilot with a marketplace returns desk. Measure auto-approval rate, fraud catches and minutes saved per claim. AWS credits fund video volume, S3 and the AgentX host.

Try it

Live demo: [demo URL]
Repository:
github.com/vnmoorthy/claimsight
Deck and storyboard in docs/

Thank you.