

ClaimSight

Refunds with receipts.

An after-sales teammate that watches the customer's evidence video, applies your refund policy, executes the refund itself, and catches returns fraud by matching footage across accounts.

CLAIM C-1042 · ORDER A1042

\$24.00 refunded

Auto-approved · 18.3 s

Evidence chip on rim @ 0:03

Policy P1 P2 P4 ✓

Transaction txn_8f3a...

EdgeOne Makers Memories.ai AgentX WorkBuddy

FRAUD TWIN DETECTED

Same footage as C-1042 from a different account (similarity 0.91). Claim frozen, manager pinged in Slack.

Every refund still needs a human to squint at a photo.

\$101B

lost to return fraud

US retailers, 2023. \$13.70 of every \$100 returned is fraud or abuse. (NRF and Appriss Retail, 2023 Consumer Returns in the Retail Industry)

0

people watching the evidence

Customers send photos and videos. Agents skim them, guess, and either overpay or escalate everything to a queue.

3

claims, one cracked mug

The same damaged item gets filmed for three claims from three "customers". No refund queue can see across accounts.

An after-sales teammate that finishes the job.

1

Sees the evidence

Memories.ai indexes the customer's 10-second clip: "white ceramic mug, chip on the rim at 0:03". Captions, frames and a summary become the claim's evidence record.

2

Applies your policy

Orders and policy clauses P1–P6 live as structured JSON in Makers storage. No vector database. Every decision cites the clauses it used.

3

Executes the transaction

Refund or replacement is written to the ledger by a cloud function that re-checks the limits itself. Humans are pulled in only above \$75 or on a fraud signal.

CLAIM C-1042 · ORDER A1042

\$24.00 refunded

Auto-approved · 18.3 s

Evidence	chip on rim @ 0:03
Policy	P1 P2 P4 ✓
Transaction	txn_8f3a...
EdgeOne Makers Memories.ai AgentX WorkBuddy	

FRAUD TWIN DETECTED
Same footage as C-1042 from a different account (similarity 0.91). Claim frozen, manager pinged in Slack.

One loop, from evidence to ledger entry.

DATA

Video + order

A 10-second clip.
Orders and policy as
JSON in Makers
storage.

INTELLIGENCE

See the damage

Memories.ai captions
the clip and image-
searches every prior
claim for a twin.

EXPERIENCE

Decision Card

Streaming chat shows
the reasoning, cited
clauses, evidence line
and txn id.

EXECUTION

Policy, twice

The agent applies it;
the refund service
enforces it again.
Over limit? A human
in Slack.

TRANSACTION

Refund issued

Ledger entry, decision
record, AgentX trace,
dashboard row.

Most teams stop at a chat answer. ClaimSight reaches the last layer: a transaction with a receipt.

A chipped mug. A refund. A fraud catch.

0:00 Claim

"My mug arrived chipped, order A1042." Evidence clip attached. The trace lights up: order, policy, evidence, fraud, execute, record.

0:30 Decision

Damage seen at 0:03. Clauses P1, P2, P4 satisfied. \$24 is under the \$75 limit. Refund executed, transaction id on the card.

1:00 Twin

Same footage from a different account. Image search finds it at 0.93 similarity. Claim frozen, escalated to Slack, denied by the manager in the Refund Desk.

The screenshot displays the ClaimSight interface, which is a dark-themed dashboard for managing claims. At the top, there's a header with 'ClaimSight' and 'refunds with receipts'. Below this, a navigation bar includes 'Chat', 'Refund Desk', and a 'DEMO DATA' button. The main content area is divided into several sections:


- Order A1042**: A summary card showing 'Chipped mug - A1042 (clean auto-approve)' and a description: 'My ceramic mug arrived with a chip on the rim - I'd like a refund.'
- ClaimSight 08:13 AM**: A chat window showing a message from a customer: 'Thanks for the video - I watched it and could see the chip and crack at 0:02 on the white ceramic mug. Order A1042 was delivered 3 days ago, inside our 30-day return window (P1: Refund window is 30 days from delivery). The damage is clearly visible in your evidence (P2: Damage must be visible in customer evidence). We normally offer a replacement first for kitchen items (P4: Kitchen and lighting items: offer replacement first when in stock), but the white ceramic mug is currently out of stock, so I've refunded you instead. I've issued a refund of \$24.00 to your original payment method - transaction txn_mu8j28hre418908d. It should appear within 3-5 business days.'
- CLAIMSIGHT - DECISION**: A card showing the decision for 'C-A1042-8103'. It includes a green 'Refund issued' button and a '\$24.00' amount. Below this, a table lists details: CUSTOMER (Alice Moreno, a_alice), EVIDENCE (White ceramic mug, chip and crack at 0:02), POLICY CLAUSES (P2, P4), and TRANSACTION (txn_mu8j28hre418908d).
- FRAUD CHECK**: A section showing 'No matching footage in other claims' and a latency of 875 ms.
- Decision trace**: A section showing 'Decided in 810 ms' and a 'POLICY ENGINE' link.
- EVIDENCE**: A section showing a video clip titled 'Chipped mug - A1042 (clean auto-approve)' with a thumbnail image of a mug.
- PIPELINE**: A section showing a list of steps: Order (lookup_order, <1 ms), Policy (get_policy, <1 ms), Evidence (inspect_evidence, <1 ms), Fraud (fraud_check, <1 ms), Execute (execute_refund, 89 ms), and Record.

At the bottom, there's a 'SCENARIOS' section with tabs for 'A1042 Chipped mug', 'A1043 Same mug, other account', and 'A1050 Cracked headphones'. Below this, there's a form for 'ORDER A1042' with a description 'White ceramic mug - \$24.00 - delivered 3d ago' and buttons for 'Attach evidence' and 'Demo clip'. A text input field for 'Describe the problem with your order...' and a 'File claim' button are also present.

Agent runtime, cloud functions and storage on EdgeOne Makers; co-hosts do the seeing, gating and reviewing.

ClaimSight architecture

Evidence video → policy decision → executed transaction, with every step traced and gated




The same mug, a different account.

Every evidence clip is indexed into the claims collection. Before paying, ClaimSight takes a frame from the new clip and runs an image search across every prior claim.

0.93

similarity to claim C-1042, filed by another customer

- Policy P5 freezes the claim instead of paying it.
- The manager gets a Slack ping through the WorkBuddy Refund Desk skill.
- One click in the desk: approve or deny, with a note that lands in the audit record.



Shipped like software.

30

golden claims

clean, over-limit, out-of-window, replacement-first, fraud twin, non-returnable

3

online monitors

PII, secrets and prompt-injection scorers on every reply

7.5

gate floor out of 10

plus a no-regression check against the previous run; CI fails the build below it

1

trace per claim

every tool call, latency and decision in AgentX, linked to the eval run

The screenshot shows the 'AGENTX Observe - Live Traces' interface. It features a sidebar with navigation options like Overview, Monitor, Review, Scors, Topics, Evaluate, Insights, and Playground. The main area displays a table of live traces. The table has columns for TIME, TRACE, SOURCE, TOOL QUALITY, SCORE, IN TOKENS, OUT TOKENS, and LATENCY. The traces are for 'agent claimsight.claim' and show a 'pass' status with a score of 4 or 5. The latency is consistently low, ranging from 1ms to 4ms. The interface also includes a search bar and a 'Framework' dropdown.

TIME	TRACE	SOURCE	TOOL QUALITY	SCORE	IN TOKENS	OUT TOKENS	LATENCY
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 4	70.0000	-	-	1ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 4	70.0000	-	-	2ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 4	70.0000	-	-	1ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	4ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	2ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	2ms
6:20 AM	agent claimsight.claim	agent claimsight.claim	pass 5	70.0000	-	-	3ms

AgentX Observe: one trace per ClaimSight claim, straight from the running backend.

Limits live in the refund service, not the prompt: the function re-checks amount, window and approval header, so the model cannot exceed policy even when asked to. The agent has no shell, no browser and no keys. If the gateway blinks, a deterministic policy engine runs the same steps and produces the same card.

Every sponsor tool does real work.

EdgeOne Makers

Agent runtime (session mode), custom MCP tools, Blob storage, AI Gateway models, streaming UI, tracing, one-command deploy.

Memories.ai

Video Datalake: upload, index, summary and captions, frame extraction, image search across the claims collection for the fraud twin.

AgentX

Self-hosted trace-eval: golden dataset, policy judge, deploy gate with no-regression check, online PII and secrets monitors.

WorkBuddy

Refund Desk skill: pulls escalations, summarizes evidence and fraud matches, takes approve or deny, posts to Slack; daily .docx report.

VeloDB

claims_events table and dashboard SQL: auto-approval by hour, refunded dollars, fraud clusters, p95 latency.

AWS

S3 evidence archive on every upload (private, versioned, encrypted) and a CloudFormation stack for the EC2 host that runs AgentX and the live-evidence relay.

Data → Intelligence → Experience → Execution → Transaction. **One loop.**

Why it wins

Returns fraud is a \$101B problem with a named buyer on every marketplace. The demo reaches a transaction, and the fraud catch is a moment judges remember.

Next 30 days

One pilot with a marketplace returns desk. Measure auto-approval rate, fraud catches and minutes saved per claim. AWS credits fund video volume, S3 and the AgentX host.

Try it

Live demo: [demo URL]
Repository:
github.com/vnmoorthy/claimsight
Deck and storyboard in docs/

Thank you.